

Scheduling

- Minimizing time in the system
- Scheduling with deadlines

Minimizing time in the system


- Assume we submit n jobs into a system at the same time.
 - The service time of job i , t_i , known in advance
- Problem
 - Design an algorithm that minimizes the average response time
 - This is equivalent to
 - Minimizing the total response time
 - $T = \sum_{i=1}^n (\text{response time of job } i)$

A greedy algorithm

- Algorithm
 1. Sort the jobs by their service times
 2. Repeat
 1. Serve the job with minimal service time among the remaining jobs
- Analysis
 - Step 1: $O(n \log n)$
 - Step 2: $\Theta(n)$
 - Total: $O(n \log n)$

Optimality of the greedy scheduling algorithm

- Theorem 6.5.1. The greedy algorithm is optimal



Compares schedules P and P'. Job a at P' leaves at the same time as job b in P. Jobs b and a+1 to b-1 in P' leave earlier than the corresponding jobs in P.

Optimality of the greedy scheduling algorithm

$$\begin{aligned} T(P) &= s_1 + (s_1 + s_2) + \dots (s_1 + s_2 + \dots s_n) \\ &= ns_1 + (n-1)s_2 + \dots + 1s_n \\ &= \sum_{k=1}^n (n-k+1)s_k \end{aligned}$$

$$T(P) = (n-a+1)s_a + (n-b+1)s_b + \sum_{k=1, k \neq a, b}^n (n-k+1)s_k$$

$$T(P') = (n-a+1)s_b + (n-b+1)s_a + \sum_{k=1, k \neq a, b}^n (n-k+1)s_k$$

$$T(P) - T(P') = (n-a+1)(s_a - s_b) + (n-b+1)(s_b - s_a) = (b-a)(s_a - s_b) > 0$$

Scheduling with deadlines

- We have a set of jobs to execute, each of which take unit time.
- No concurrency and parallel: at any time we can execute exactly one job
- Job i earn us a profit $g_i > 0$ if and only if it is executed no later than time d_i
- Problem: find a scheduling that maximize the total profits.

Example

i	1	2	3	4
g_i	50	10	15	30
d_i	2	1	2	1

Sequence	Profit	Sequence	Profit
1	50	2, 1	60
2	10	2, 3	25
3	15	3, 1	65
4	30	4, 1	80
1, 3	65	4, 3	45

← Optimum

A greedy algorithm


- Feasible
 - A set of jobs is feasible if there exists at least one sequence (also called feasible) that allows all jobs in the set to be executed no late than their respective deadlines.
- Greedy algorithm
 - At each step, add the job with the highest profit (g_i) among those not yet considered, providing that the chosen set of jobs remains feasible.
- Theorem: the greedy algorithm above is optimal

Proof of the optimality of the greedy algorithm

- Suppose the greedy algorithm choose the set of jobs I , and the set J is optimal. Let S_I and S_J be feasible sequences, possibly including gaps
 - Rearrange the jobs in S_I and S_J , we can obtain two feasible sequences S_I' and S_J' , such that every job common to I and J is scheduled at the same time in both sequences
 - Compare each pair of slots at the same time and show that I' and J' yield the same profit


Rearranging Jobs, case I

- Job a occurs in both feasible sequences and is scheduled at time t_i and t_j respectively
 - $t_i = t_j$, we do not need to do anything



Rearranging Jobs, case II

- Job a occurs in both feasible sequences and is scheduled at time t_i and t_j respectively
 - $t_i < t_j$, we will try to move job a at t_i to the slot t_j
 - Slot t_j is a gap. Simply move
 - Job b is at slot t_j . Exchange job a and b .



Rearranging Jobs, case III

- Job a occurs in both feasible sequences and is scheduled at time t_i and t_j respectively
 - $t_i > t_j$, similar to case II.

Compare jobs in the rearranged sequences

- If some job a is schedule in S_1' opposite a gap in S_j' , a does not belong J . $J \cup \{a\}$ is a feasible sequence and more profitable.
- If some job b is scheduled in S_j' opposite a gap in S_1' . The set $I \cup \{b\}$ will be feasible and the greedy algorithm would have included b in I .
- The only remaining possibility is that some job a is schedules in S_1' opposite a different job b in S_j' .
 - $g_a > g_b$. J is not optimal
 - $g_a < g_b$. The greedy algorithm should have chosen b first.
 - The only remaining possibility is $g_a = g_b$.

Lemma

- Let J be a set of jobs. Suppose without loss of generality that the jobs are numbered so that $d_1 \leq d_2 \leq \dots \leq d_k$
- Then the set J is feasible if and only if the sequence 1, 2, ..., k is feasible
 - Key to the proof: the pigeonhole principle

An implementation

Sequence($d[n+1]$)
 {
 // assuming the profit is sorted in decreasing order
 $d[0] = s[0] = 0$; // sentinels
 $s[1] = 1$; // the first job is scheduled
 $k = 1$; // # jobs already scheduled

for ($i=2$; $i \leq n$; $i++$) { // try job i
 // invariant: $d[s[0]] \leq d[s[1]] \leq \dots \leq d[s[k]]$
 $r = k$;
 while ($d[s[r]] > \max(d[i], r)$) $r--$; // $r+1$ is the earliest slot for job i
 if ($d[i] > r$) { // feasible if insert i
 for ($j=k$; $j > r$; $j--$)
 $s[j+1] = s[j]$;
 $s[r+1] = i$;
 $k++$;
 }
 }
 return s ;
 }

Worst Case: $O(n^2)$

Example

i	1	2	3	4	5	6
g_i	20	15	10	7	5	3
d_i	3	1	1	3	1	3

Initialization: $d[s[i]]$

3						
1						

Try 2

1	3					
2	1					

Try 3 unchanged

1	3	3				
2	1	4				

Try 4

Try 5, 6 unchanged

Optimal sequence: 2, 1, 4


A more efficient algorithm begins with the lemma below

- Lemma 6.6.4
 - A set of jobs J is feasible if and only if we can construct a feasible sequence including all the jobs in J as follows. Start with an empty schedule of length n . Then for each job $i \in J$ in turn, schedule i at time t , where t is the largest integer such that $1 \leq t \leq \min(n, d_i)$ and the job to be executed at time t is not yet decided.
 - The “if” is obvious
 - The “only if”

Proof of “only of”

- If the sequence is feasible, then it can be scheduled in n slots.
- When we try to add a new job, the sequence being built always contains a gap
- Assume by contradiction that we are unable to add a job of deadline d
 - s is the next available gap with $s > d$
 - We can show that all scheduled jobs in slot 1 to $s-1$ have deadlines $\leq s-1$, which results in s jobs with deadline $\leq s-1$ and thus no schedule is feasible

Use disjoint set to construct feasible sequence



$$n_i = \max\{k \leq t \mid \text{position } k \text{ is free}\}$$

Two positions i and j are in the same set if $n_i = n_j$

$F(K)$ is the smallest member of set K

Algorithm

```
Sequence2(d[n+1])
{
    // assuming the profit is sorted
    // in decreasing order
    for (i=0; i<=n; i++) {
        s[i] = 0;
        F[i] = i;
        initialize set i;
    }

    // greedy loop
    for (i=1; i<=n; i++) { // try job i
        k = find(min(n, d[i]));
        m = F[k];
        if (m != 0) {
            s[m] = i;
            l = find(m-1);
            F[k] = F[l];
            merge(k, l);
        }
    }


    // compress the solution
    k=0
    for (i=1; i<=n; i++) {
        if (s[i] > 0) {
            k = k+1;
            s[k] = s[i];
        }
    }
    return s;
}
```

Analysis: $2 \cdot n$ find and n merge
 Cost: $O(n \alpha(2n, n))$
 Sorting cost for the assumption.

Example

i	1	2	3	4	5	6
g_i	20	15	10	7	5	3
d_i	3	1	1	3	1	3


Initialization: $l = \min(6, \max(d_i)) = 3$




Example

i	1	2	3	4	5	6
g_i	20	15	10	7	5	3
d_i	3	1	1	3	1	3

Try 1: $d_1=3$, assign task 1 to position 3



Try 2: $d_2=1$, assign task 2 to position 1




Example

i	1	2	3	4	5	6
g_i	20	15	10	7	5	3
d_i	3	1	1	3	1	3

Try 3: $d_3=1$, no free slot is available

Try 4: $d_4=3$, assign task 4 to position 2



Try 5 and try 6, no free position is available