

Constraints and Triggers

Slides adapted from <http://infolab.stanford.edu/~ullman/fcdb.html>

1

Constraints

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
- Triggers are only executed when a specified condition occurs, e.g., insertion of a tuple.
 - Easier to implement than complex constraints.

2

Kinds of Constraints

- **Keys**
- **Foreign-key**, or referential-integrity
- **Value-based** constraints
 - Constrain values of a particular attribute
- **Tuple-based** constraints
 - Relationship among components
- **Assertions**: any SQL boolean expression

3

Review: Single-Attribute Keys

- Place PRIMARY KEY or UNIQUE NOT NULL after the type in the declaration of the attribute.
- **Example:**

```
CREATE TABLE Beers (  
    name CHAR(20) UNIQUE NOT NULL,  
    manf CHAR(20)  
);
```

4

Review: Multiattribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar        CHAR(20),  
    beer       VARCHAR(20),  
    price      REAL,  
    PRIMARY KEY (bar, beer)  
);
```

5

Foreign Keys

- Values appearing in attributes of one relation must appear together in certain attributes of another relation.
- **Example:** in **Sells(bar, beer, price)**, we expect that a beer value also appears in Beers table as Beers.name .

6

Expressing Foreign Keys

- Use keyword REFERENCES, either:
 1. After an attribute (for one-attribute keys).
 2. As an element of the schema:
FOREIGN KEY (<list of attributes>)
REFERENCES <relation> (<attributes>)
- Referenced attributes must be declared PRIMARY KEY or UNIQUE.

7

Example: With Attribute

```
CREATE TABLE Beers (  
    name    CHAR(20) PRIMARY KEY,  
    manf    CHAR(20) );  
  
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20) REFERENCES Beers(name),  
    price    REAL );
```

8

Example: As Schema Element

```
CREATE TABLE Beers (  
    name    CHAR(20) PRIMARY KEY,  
    manf    CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    FOREIGN KEY (beer) REFERENCES  
        Beers (name) );
```

9

Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from relation R to relation S , two violations are possible:
 1. An insert or update to R introduces values not found in S .
 2. A deletion or update to S causes some tuples of R to “dangle.”

10

Actions Taken --- (1)

- **Example**: suppose $R = \text{Sells}$, $S = \text{Beers}$.
- An insert or update to **Sells** that introduces a nonexistent beer must be rejected.
- A deletion or update to **Beers** that removes a beer value found in some tuples of **Sells** can be handled in three ways (next slide).

11

Actions Taken --- (2)

1. **Default** : Reject the modification.
2. **Cascade** : Make the same changes in **Sells**.
 - **Deleted beer**: delete **Sells** tuple.
 - **Updated beer**: change value in **Sells**.
3. **Set NULL** : Change the beer to NULL.

12

Example: Cascade

- Delete the Bud tuple from Beers:
 - Then delete all tuples from Sells that have beer = 'Bud'.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
 - Then change all Sells tuples with beer = 'Bud' to beer = 'Budweiser'.

13

Example: Set NULL

- Delete the Bud tuple from Beers:
 - Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
 - Same change as for deletion.

14

Choosing a Policy

- When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- Follow the foreign-key declaration by:
ON [UPDATE/DELETE][SET NULL/CASCADE]
- Otherwise, the default (reject) is used.

15

Example: Setting Policy

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

16

Attribute-Based Checks

- Constraints on the value of a particular attribute.
- Add CHECK(<condition>) to the declaration for the attribute.
- The condition may use the name of the attribute, but **any other relation or attribute name must be in a subquery.**

17

Example: Attribute-Based Check

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20) CHECK (beer IN  
        (SELECT name FROM Beers)),  
    price  REAL CHECK ( price <= 5.00 )  
);
```

18

Timing of Checks

- Attribute-based checks are performed only when a value for that attribute is inserted or updated.
 - **Example:** `CHECK (price <= 5.00)` checks every new price and rejects the modification (for that tuple) if the price is more than \$5.
 - **Example:** `CHECK (beer IN (SELECT name FROM Beers))` not checked if a beer is deleted from Beers (unlike foreign-keys).

19

Tuple-Based Checks

- `CHECK (<condition>)` may be added as a relation-schema element.
- The condition may refer to any attribute of the relation.
 - But other attributes or relations require a subquery.
- Checked on insert or update only.

20

Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar        CHAR(20),  
    beer       CHAR(20),  
    price      REAL,  
    CHECK (bar = 'Joe's Bar' OR  
           price <= 5.00)  
);
```

21

Assertions

- These are database-schema elements, like relations or views.
- Defined by:

```
CREATE ASSERTION <name>  
    CHECK (<condition>);
```
- Condition may refer to any relation or attribute in the database schema.

22

Example: Assertion

- In **Sells(bar, beer, price)**, no bar may charge an average of more than \$5.

```
CREATE ASSERTION NoRipoffBars CHECK (  
  NOT EXISTS (  
    SELECT bar FROM Sells  
    GROUP BY bar  
    HAVING 5.00 < AVG(price)  
  )  
);
```

Bars with an average price above \$5

23

Example: Assertion

- In **Drinkers(name, addr, phone)** and **Bars(name, addr, license)**, there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
  (SELECT COUNT(*) FROM Bars) <=  
  (SELECT COUNT(*) FROM Drinkers)  
);
```

24

Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.
 - **Example:** No change to Beers can affect FewBar. Neither can an insertion to Drinkers.

25

Triggers: Motivation

- Assertions are powerful, but the DBMS often can't tell when they need to be checked.
- Attribute- and tuple-based checks are checked at known times, but are not powerful.
- Triggers let the user decide when to check for any condition.

26

Event-Condition-Action Rules

- Another name for “trigger” is *ECA rule*, or *event-condition-action* rule.
- *Event* : typically a type of database modification, e.g., “insert on Sells.”
- *Condition* : Any SQL boolean-valued expression.
- *Action* : Any SQL statements.

27

Preliminary Example: A Trigger

- Instead of using a foreign-key constraint and rejecting insertions into *Sells(bar, beer, price)* with unknown beers, a trigger can add that beer to *Beers*, with a NULL manufacturer.

28

Example: Trigger Definition

```
CREATE TRIGGER BeerTrig
  AFTER INSERT ON Sells
  REFERENCING NEW ROW AS NewTuple
  FOR EACH ROW
  WHEN (NewTuple.beer NOT IN
        (SELECT name FROM Beers))
  INSERT INTO Beers(name)
  VALUES(NewTuple.beer);
```

The event

The condition

The action

29

Options: CREATE TRIGGER

- CREATE TRIGGER <name>
- Or:
CREATE OR REPLACE TRIGGER <name>
 - Useful if there is a trigger with that name and you want to modify the trigger.

30

Options: The Event

- AFTER can be BEFORE.
 - Also can be INSTEAD OF, if the relation is a view.
- INSERT can be DELETE or UPDATE.
 - And UPDATE can be UPDATE . . . ON a particular attribute.

31

Options: FOR EACH ROW

- Triggers are either “row-level” or “statement-level.”
- FOR EACH ROW indicates row-level; its absence indicates statement-level.
- *Row level triggers* : execute once for each modified tuple.
- *Statement-level triggers* : execute once for a SQL statement, regardless of how many tuples are modified (0, 1, or more).

32

Options: REFERENCING

- INSERT statements imply a new tuple (for row-level) or new table (for statement-level).
 - The “table” is the set of inserted tuples.
- DELETE implies an old tuple or table.
- UPDATE implies both.
- Refer to these by
[NEW OLD][TUPLE TABLE] AS <name>

33

Options: The Condition

- Any boolean-valued condition.
- Evaluated on the database as it would exist before or after the triggering event, depending on whether BEFORE or AFTER is used.
 - But always before the changes take effect.
- Access the new/old tuple/table through the names in the REFERENCING clause.

34

Options: The Action

- There can be more than one SQL statement in the action.
 - Surround by BEGIN . . . END if there is more than one.
- But queries make no sense in an action, so we are really limited to modifications.

35

Another Example

- Using **Sells(bar, beer, price)** and a unary relation **RipoffBars(bar)**, maintain a list of bars that raise the price of any beer by more than \$1.

36

The Trigger

CREATE TRIGGER PriceTrig

AFTER UPDATE OF price ON Sells

The event –
only changes
to prices

REFERENCING

OLD ROW AS ooo

NEW ROW AS nnn

Updates let us
talk about old
and new tuples

Condition:
a raise in
price > \$1

We need to consider
each price change

FOR EACH ROW

WHEN(nnn.price > ooo.price + 1.00)

INSERT INTO RipoffBars

VALUES(nnn.bar);

When the price change
is great enough, add
the bar to RipoffBars

37