

# Relational Algebra

Slides adapted from <http://infolab.stanford.edu/~ullman/fcdb.html>



## What is a Data Model?

1. Mathematical representation of data.
  - Examples: relational model = tables; semistructured model = trees/graphs.
2. Operations on data.
3. Constraints.



2 of 42

## A Relation is a Table

Attributes (column headers)	name	manufacture
Tuples (rows)	Winterbrew	Pete's
	Bud Lite	Anheuser-Busch

Beers

Relation  
name

3 of 42

## Schemas

- **Relation schema** = relation name and attribute list.
  - Optionally: types of attributes.
  - Example: **Beers(name, manf)** or **Beers(name: string, manf: string)**
- **Database** = collection of relations.
- **Database schema** = set of all relation schemas in the database.

4 of 42

## Why Relations?

- Very simple model.
- *Often* matches how we think about data.
- Abstract model that underlies SQL, the most important database language today.

5 of 42

## A Running Example

Beers(name, manf)  
Bars(name, addr, license)  
Drinkers(name, addr, phone)  
Likes(drinker, beer)  
Sells(bar, beer, price)  
Frequents(drinker, bar)

- Underline = *key* (tuples cannot have the same value in all key attributes).
  - Excellent example of a constraint.

6 of 42

## Database Schemas in SQL

- SQL is primarily a query language, for getting information from a database.
- But SQL also includes a *data-definition* component for describing database schemas.

7 of 42

## Creating (Declaring) a Relation

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
)
```

- To delete a relation:

```
DROP TABLE <name>;
```

8 of 42

## Elements of Table Declarations

- Most basic element: an attribute and its type.
- The most common types are:
  - INT or INTEGER (synonyms).
  - REAL or FLOAT (synonyms).
  - CHAR(*n*) = fixed-length string of *n* characters.
  - VARCHAR(*n*) = variable-length string of up to *n* characters.

9 of 42

## Example: Create Table

```
CREATE TABLE Sells (
    bar      CHAR(20),
    beer     VARCHAR(20),
    price    REAL
);
```

10 of 42

## SQL Values

- Integers and reals are represented as you would expect.
- Strings are too, except they require single quotes.
  - Two single quotes = real quote, e.g., 'Joe''s Bar'.
- Any value can be NULL.

11 of 42

## Dates and Times

- DATE and TIME are types in SQL.
- The form of a date value is:  
`DATE 'yyyy-mm-dd'`
  - **Example:** DATE '2019-09-10' for Sept. 10, 2019.

12 of 42

## Times as Values

- The form of a time value is:

TIME 'hh:mm:ss'

with an optional decimal point and fractions of a second following.

- Example: TIME '11:00:02.5' = two and a half seconds after 11:00AM.

13 of 42

## Declaring Keys

- An attribute or list of attributes may be declared PRIMARY KEY or UNIQUE.
- Either says that no two tuples of the relation may agree in all the attribute(s) on the list.
  1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
  2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL.

14 of 42

## Declaring Single-Attribute Keys

- Place PRIMARY KEY after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (
    name CHAR(20) PRIMARY KEY,
    manf CHAR(20)
);
```

15of 42

## Declaring Multiattribute Keys

- A key declaration can also be another element in the list of elements of a CREATE TABLE statement.
- This form is essential if the key consists of more than one attribute.
  - May be used even for one-attribute keys.

16of 42

## Example: Multiattribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (
    bar      CHAR(20),
    beer     VARCHAR(20),
    price    REAL,
    PRIMARY KEY (bar, beer)
);
```

17 of 42

## What is an “Algebra”

- Mathematical system consisting of:
  - **Operands** - variables or values from which new values can be constructed.
  - **Operators** - symbols denoting procedures that construct new values from given values.

18 of 42

## What is Relational Algebra?

- An algebra whose operands are relations or variables that represent relations.
- Operators are designed to do the most common things that we need to do with relations in a database.
  - The result is an algebra that can be used as a *query language* for relations.

19 of 42

## Core Relational Algebra

- Union, intersection, and difference.
  - Usual set operations, but *both operands must have the same relation schema*.
- Selection: picking certain rows.
- Projection: picking certain columns.
- Products and joins: compositions of relations.
- Renaming of relations and attributes.

20 of 42

## Selection

- $R1 := \sigma_C(R2)$ 
  - $C$  is a condition (as in “if” statements) that refers to attributes of  $R2$ .
  - $R1$  is all those tuples of  $R2$  that satisfy  $C$ .

21 of 42

## Example: Selection

Relation Sells:

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

JoeMenu :=  $\sigma_{\text{bar}=\text{"Joe's"}}(\text{Sells})$ :

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75

22 of 42

## Projection

*to select rows / tuples / records*

- $R1 := \Pi_L(R2)$

- $L$  is a list of attributes from the schema of  $R2$ .
- $R1$  is constructed by looking at each tuple of  $R2$ , extracting the attributes on list  $L$ , in the order specified, and creating from those components a tuple for  $R1$ .
- Eliminate duplicate tuples, if any.

23 of 42

## Example: Projection

*to select column(s) / attributes*

Relation Sells:

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Prices :=  $\Pi_{beer, price}(Sells)$ :

beer	price
Bud	2.50
Miller	2.75
Miller	3.00

24 of 42

## Product

- $R3 := R1 \times R2$

- Pair each tuple  $t_1$  of  $R1$  with each tuple  $t_2$  of  $R2$ .
- Concatenation  $t_1t_2$  is a tuple of  $R3$ .
- Schema of  $R3$  is the attributes of  $R1$  and then  $R2$ , in order.
- But beware attribute  $B$  of the same name in  $R1$  and  $R2$ : use  $R1.B$  and  $R2.B$ .

25 of 42

### Example: $R3 := R1 \times R2$

R1(	A,	B)	)	2 tuples	R3(	A,	R1.B,	R2.B,	C	)
	1	2				1	2	5	6	
	3	4				1	2	7	8	

R2(	B,	C)	)	3 tuples	R3(	A,	R1.B,	R2.B,	C	)
	5	6				1	2	9	10	
	7	8				3	4	5	6	
	9	10				3	4	7	8	

$2 \times 3 = 6$  tuples

26 of 42

## Theta-Join

$\bowtie_c$  : apply  $\times$ , then  
apply condition  $\sigma_c$

- $R3 := R1 \bowtie_c R2$ 
  - Take the product  $R1 \times R2$ .
  - Then apply  $\sigma_c$  to the result.
- As for  $\sigma$ ,  $C$  can be any boolean-valued condition.
  - Historic versions of this operator allowed only  $A \theta B$ , where  $\theta$  is  $=$ ,  $<$ , etc.; hence the name “theta-join.”

27 of 42

## Example: Theta Join

Sells(	bar,	beer,	price	)	Joe	Mrs.	Bars(	name, addr	)
	Joe's	Bud	2.50		Sue	R st			
	Joe's	Miller	2.75		Joe	M st			
	Sue's	Bud	2.50		Sue	R st			
	Sue's	Coors	3.00		Joe	M st			
					Sue	R st			

BarInfo := Sells  $\bowtie_{Sells.bar = Bars.name}$  Bars

BarInfo(	bar,	beer,	price,	name,	addr	)
	Joe's	Bud	2.50	Joe's	Maple St.	
	Joe's	Miller	2.75	Joe's	Maple St.	
	Sue's	Bud	2.50	Sue's	River Rd.	
	Sue's	Coors	3.00	Sue's	River Rd.	

28 of 42

Joe

## Natural Join

- A useful join variant (*natural* join) connects two relations by:
  - Equating attributes of the same name, and
  - Projecting out one copy of each pair of equated attributes.
- Denoted  $R_3 := R_1 \bowtie R_2$ .

29 of 42

## Example: Natural Join

Sells( bar, beer, price )		
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Coors	3.00

Bars( bar, addr )	
Joe's	Maple St.
Sue's	River Rd.

BarInfo := Sells  $\bowtie$  Bars

Note: *Bars.name* has become *Bars.bar* to make the natural join "work."

BarInfo( bar, beer, price, addr )			
Joe's	Bud	2.50	Maple St.
Joe's	Miller	2.75	Maple St.
Sue's	Bud	2.50	River Rd.
Sue's	Coors	3.00	River Rd.

30 of 42

## Renaming

- The  $\rho$  operator gives a new schema to a relation.
- $R1 := \rho_{R1(A1, \dots, An)}(R2)$  makes  $R1$  be a relation with attributes  $A1, \dots, An$  and the same tuples as  $R2$ .
- Simplified notation:  $R1(A1, \dots, An) := R2$ .

31 of 42

take attributes  
 $A_1, A_2, \dots, A_n$  of  
 $R_1$  to make a  
relation  $R_2$

## Example: Renaming

Bars( name, addr )	
Joe's	Maple St.
Sue's	River Rd.

$R(\text{bar}, \text{addr}) := \text{Bars}$

R( bar, addr )	
Joe's	Maple St.
Sue's	River Rd.

to take attributes 'bar'  
and 'addr' of  $R$   
to create relation  
'Bars'

32 of 42

## Building Complex Expressions

- Combine operators with parentheses and precedence rules.
- Three notations, just as in arithmetic:
  1. Sequences of assignment statements.
  2. Expressions with several operators.
  3. Expression trees.

33 of 42

## Sequences of Assignments

- Create temporary relation names.
- Renaming can be implied by giving relations a list of attributes.
- Example:  $R3 := R1 \bowtie_c R2$  can be written:

$R4 := R1 \times R2$

$R3 := \sigma_c(R4)$

34 of 42

## Expressions in a Single Assignment

- Example: the theta-join  $R3 := R1 \bowtie_c R2$  can be written:  $R3 := \sigma_c(R1 \times R2)$
- Precedence of relational operators:
  1.  $[\sigma, \pi, \rho]$  (highest)
  2.  $[x, \bowtie]$
  3.  $\cap$
  4.  $[\cup, -]$

35 of 42

## Expression Trees

- Leaves are operands - either variables standing for relations or particular, constant relations.
- Interior nodes are operators, applied to their child or children.

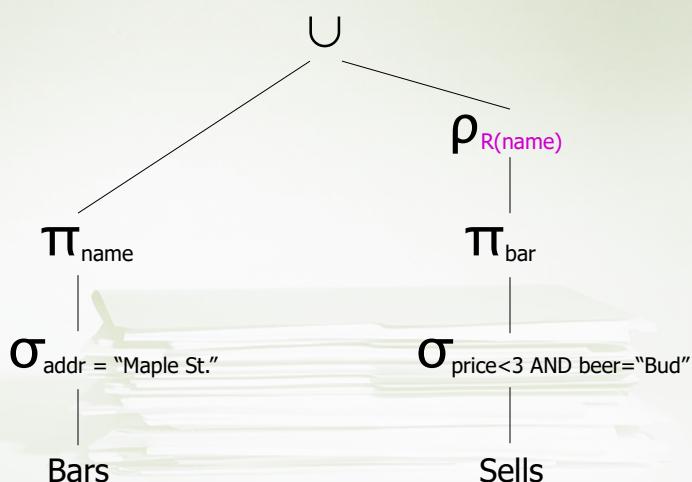
36 of 42

## Example: Tree for a Query

- Using the relations  $\text{Bars}(\text{name}, \text{addr})$  and  $\text{Sells}(\text{bar}, \text{beer}, \text{price})$ , find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.

37 of 42

As a Tree:



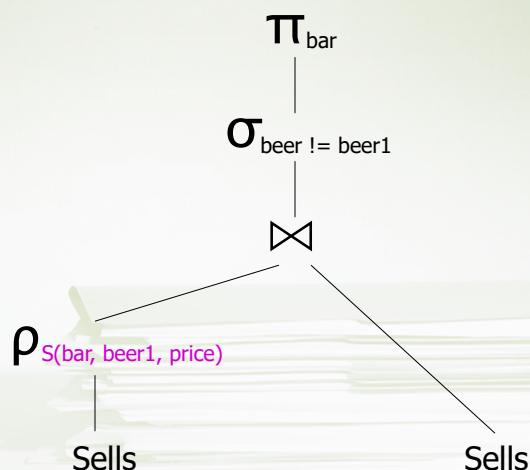
38 of 42

## Example: Self-Join

- Using  $\text{Sells}(\text{bar}, \text{beer}, \text{price})$ , find the bars that sell two different beers at the same price.
- **Strategy:** by renaming, define a copy of Sells, called  $\text{S}(\text{bar}, \text{beer1}, \text{price})$ . The natural join of Sells and S consists of quadruples  $(\text{bar}, \text{beer}, \text{beer1}, \text{price})$  such that the bar sells both beers at this price.

39 of 42

## The Tree



40 of 42

## Schemas for Results

- **Union, intersection, and difference:** the schemas of the two operands must be the same, so use that schema for the result.
- **Selection:** schema of the result is the same as the schema of the operand.
- **Projection:** list of attributes tells us the schema.

41 of 42

## Schemas for Results --- (2)

- **Product:** schema is the attributes of both relations.
  - Use  $R1.A$  and  $R2.A$ , etc., to distinguish two attributes named A.
- **Theta-join:** same as product.
- **Natural join:** union of the attributes of the two relations.
- **Renaming:** the operator tells the schema.

42 of 42