

## Join GitHub today

GitHub is home to over 31 million developers working together to host and review code, manage projects, and build software together.

[Sign up](#)

Branch: master ▾ DB2 / src / main / java / index / BPlusTree.java

[Find file](#)[Copy path](#) dreamlegends init add all files

2c520b8 7 days ago

1 contributor

395 lines (358 sloc) 14.3 KB

[Raw](#)[Blame](#)[History](#)

```
1 package index;
2
3 import java.nio.ByteBuffer;
4 import java.util.ArrayList;
5 import java.util.Iterator;
6 import java.util.List;
7 import java.util.NoSuchElementException;
8 import java.util.Optional;
9
10 import common.Pair;
11 import databox.DataBox;
12 import databox.Type;
13 import io.Page;
14 import io.PageAllocator;
15 import table.RecordId;
16
17 /**
18  * A persistent B+ tree.
19  * <p>
20  * // Create an order 2, integer-valued B+ tree that is persisted in tree.txt.
21  * BPlusTree tree = new BPlusTree("tree.txt", Type.intType(), 2);
22  * <p>
23  * // Insert some values into the tree.
24  * tree.put(new IntDataBox(0), new RecordId(0, (short) 0));
25  * tree.put(new IntDataBox(1), new RecordId(1, (short) 1));
26  * tree.put(new IntDataBox(2), new RecordId(2, (short) 2));
27  * <p>
28  * // Get some values out of the tree.
29  * tree.get(new IntDataBox(0)); // Optional.of(RecordId(0, 0))
30  * tree.get(new IntDataBox(1)); // Optional.of(RecordId(1, 1))
31  * tree.get(new IntDataBox(2)); // Optional.of(RecordId(2, 2))
32  * tree.get(new IntDataBox(3)); // Optional.empty();
33  * <p>
34  * // Iterate over the record ids in the tree.
35  * tree.scanEqual(new IntDataBox(2)); // [(2, 2)]
36  * tree.scanAll(); // [(0, 0), (1, 1), (2, 2)]
37  * tree.scanGreaterEqual(new IntDataBox(1)); // [(1, 1), (2, 2)]
38  * <p>
39  * // Remove some elements from the tree.
40  * tree.get(new IntDataBox(0)); // Optional.of(RecordId(0, 0))
41  * tree.remove(new IntDataBox(0));
42  * tree.get(new IntDataBox(0)); // Optional.empty()
43  * <p>
44  * // Load the tree from disk.
```

```

21 * BPlusTree fromDisk = new BPlusTree("tree.txt");
22 * <p>
23 * // All the values are still there.
24 * fromDisk.get(new IntDataBox(0)); // Optional.empty()
25 * fromDisk.get(new IntDataBox(1)); // Optional.of(RecordId(1, 1))
26 * fromDisk.get(new IntDataBox(2)); // Optional.of(RecordId(2, 2))
27 */
28 public class BPlusTree {
29     public static final String FILENAME_PREFIX = "db";
30     public static final String FILENAME_EXTENSION = ".index";
31
32     private BPlusTreeMetadata metadata;
33     private Page headerPage;
34     private BPlusNode root;
35
36     // Constructors //////////////////////////////////////
37
38     /**
39     * Construct a new B+ tree which is serialized into the file `filename`,
40     * stores keys of type `keySchema`, and has order `order`. For example,
41     * `new BPlusTree("tree.txt", Type.intType(), 2)` constructs a B+ tree that
42     * is serialized to "tree.txt", that maps integers to record ids, and that
43     * has order 2.
44     * <p>
45     * If the specified order is so large that a single node cannot fit on a
46     * single page, then a BPlusTree exception is thrown. If you want to have
47     * maximally full B+ tree nodes, then use the BPlusTree.maxOrder function
48     * to get the appropriate order.
49     * <p>
50     * We reserve the first page (i.e. page number 0) of the file for a header
51     * page which contains:
52     * <p>
53     * - the key schema of the tree,
54     * - the order of the tree, and
55     * - the page number of the root of the tree.
56     * <p>
57     * All other pages are serializations of inner and leaf nodes. See
58     * writeHeader for details.
59     */
60     public BPlusTree(String filename, Type keySchema, int order)
61         throws BPlusTreeException {
62         // Sanity checks.
63         if (order < 0) {
64             String msg = String.format(
65                 "You cannot construct a B+ tree with negative order %d.",
66                 order);
67             throw new BPlusTreeException(msg);
68         }
69
70         int maxOrder = BPlusTree.maxOrder(Page.pageSize, keySchema);
71         if (order > maxOrder) {
72             String msg = String.format(
73                 "You cannot construct a B+ tree with order %d greater than the " +
74                 "max order %d.",
75                 order, maxOrder);
76             throw new BPlusTreeException(msg);
77         }
78
79         // Initialize the page allocator.
80         PageAllocator allocator = new PageAllocator(filename, true /* wipe */);
81         this.metadata = new BPlusTreeMetadata(allocator, keySchema, order);
82
83         // Allocate the header page.
84         int headerPageNum = allocator.allocPage();
85         assert (headerPageNum == 0);
86         this.headerPage = allocator.fetchPage(headerPageNum);
87
88         // Construct the root.
89         List<DataBox> keys = new ArrayList<>();

```

```

114     List<RecordId> rids = new ArrayList<>();
115     Optional<Integer> rightSibling = Optional.empty();
116     this.root = new LeafNode(this.metadata, keys, rids, rightSibling);
117
118     // Initialize the header page.
119     writeHeader(headerPage.getByteBuffer());
120 }
121
122 /**
123  * Read a B+ tree that was previously serialized to filename.
124  */
125 public BPlusTree(String filename) {
126     // Initialize the page allocator and fetch the header page.
127     PageAllocator allocator = new PageAllocator(filename, false /* wipe */);
128     Page headerPage = allocator.fetchPage(0);
129     ByteBuffer buf = headerPage.getByteBuffer();
130
131     // Read the contents of the header page. See writeHeader for information
132     // on exactly what is written to the header page.
133     Type keySchema = Type.fromBytes(buf);
134     int order = buf.getInt();
135     int rootPageNum = buf.getInt();
136
137     // Initialize members.
138     this.metadata = new BPlusTreeMetadata(allocator, keySchema, order);
139     this.headerPage = allocator.fetchPage(0);
140     this.root = BPlusNode.fromBytes(this.metadata, rootPageNum);
141 }
142
143 // Core API //////////////////////////////////////
144
145 /**
146  * Returns the value associated with `key`.
147  * <p>
148  * // Create a B+ tree and insert a single value into it.
149  * BPlusTree tree = new BPlusTree("t.txt", Type.intType(), 4);
150  * DataBox key = new IntDataBox(42);
151  * RecordId rid = new RecordId(0, (short) 0);
152  * tree.put(key, rid);
153  * <p>
154  * // Get the value we put and also try to get a value we never put.
155  * tree.get(key); // Optional.of(rid)
156  * tree.get(new IntDataBox(100)); // Optional.empty()
157  */
158 public Optional<RecordId> get(DataBox key) {
159     typecheck(key);
160     throw new UnsupportedOperationException("TODO: implement");
161 }
162
163 /**
164  * scanEqual(k) is equivalent to get(k) except that it returns an iterator
165  * instead of an Optional. That is, if get(k) returns Optional.empty(),
166  * then scanEqual(k) returns an empty iterator. If get(k) returns
167  * Optional.of(rid) for some rid, then scanEqual(k) returns an iterator
168  * * over rid.
169  */
170 public Iterator<RecordId> scanEqual(DataBox key) {
171     typecheck(key);
172     Optional<RecordId> rid = get(key);
173     if (rid.isPresent()) {
174         ArrayList<RecordId> l = new ArrayList<>();
175         l.add(rid.get());
176         return l.iterator();
177     } else {
178         return new ArrayList<RecordId>().iterator();
179     }
180 }
181
182 /**

```

```

187 * Returns an iterator over all the RecordIds stored in the B+ tree in
188 * ascending order of their corresponding keys.
189 * <p>
190 * // Create a B+ tree and insert some values into it.
191 * BPlusTree tree = new BPlusTree("t.txt", Type.intType(), 4);
192 * tree.put(new IntDataBox(2), new RecordId(2, (short) 2));
193 * tree.put(new IntDataBox(5), new RecordId(5, (short) 5));
194 * tree.put(new IntDataBox(4), new RecordId(4, (short) 4));
195 * tree.put(new IntDataBox(1), new RecordId(1, (short) 1));
196 * tree.put(new IntDataBox(3), new RecordId(3, (short) 3));
197 * <p>
198 * Iterator<RecordId> iter = tree.scanAll();
199 * iter.next(); // RecordId(1, 1)
200 * iter.next(); // RecordId(2, 2)
201 * iter.next(); // RecordId(3, 3)
202 * iter.next(); // RecordId(4, 4)
203 * iter.next(); // RecordId(5, 5)
204 * iter.next(); // NoSuchElementException
205 * <p>
206 * Note that you CAN NOT materialize all record ids in memory and then
207 * return an iterator over them. Your iterator must lazily scan over the
208 * leaves of the B+ tree. Solutions that materialize all record ids in
209 * memory will receive 0 points.
210 */
211 public Iterator<RecordId> scanAll() {
212     throw new UnsupportedOperationException("TODO: implement");
213 }
214
215 /**
216 * Returns an iterator over all the RecordIds stored in the B+ tree that
217 * are greater than or equal to 'key'. RecordIds are returned in ascending
218 * of their corresponding keys.
219 * <p>
220 * // Create a B+ tree and insert some values into it.
221 * BPlusTree tree = new BPlusTree("t.txt", Type.intType(), 4);
222 * tree.put(new IntDataBox(2), new RecordId(2, (short) 2));
223 * tree.put(new IntDataBox(5), new RecordId(5, (short) 5));
224 * tree.put(new IntDataBox(4), new RecordId(4, (short) 4));
225 * tree.put(new IntDataBox(1), new RecordId(1, (short) 1));
226 * tree.put(new IntDataBox(3), new RecordId(3, (short) 3));
227 * <p>
228 * Iterator<RecordId> iter = tree.scanGreaterEqual(new IntDataBox(3));
229 * iter.next(); // RecordId(3, 3)
230 * iter.next(); // RecordId(4, 4)
231 * iter.next(); // RecordId(5, 5)
232 * iter.next(); // NoSuchElementException
233 * <p>
234 * Note that you CAN NOT materialize all record ids in memory and then
235 * return an iterator over them. Your iterator must lazily scan over the
236 * leaves of the B+ tree. Solutions that materialize all record ids in
237 * memory will receive 0 points.
238 */
239 public Iterator<RecordId> scanGreaterEqual(DataBox key) {
240     throw new UnsupportedOperationException("TODO: implement");
241 }
242
243 /**
244 * Inserts a (key, rid) pair into a B+ tree. If the key already exists in
245 * the B+ tree, then the pair is not inserted and an exception is raised.
246 * <p>
247 * BPlusTree tree = new BPlusTree("t.txt", Type.intType(), 4);
248 * DataBox key = new IntDataBox(42);
249 * RecordId rid = new RecordId(42, (short) 42);
250 * tree.put(key, rid); // Success :)
251 * tree.put(key, rid); // BPlusTreeException :(
252 */
253 public void put(DataBox key, RecordId rid) throws BPlusTreeException {
254     typecheck(key);
255     throw new UnsupportedOperationException("TODO: implement");
256 }

```

```

252 }
253
254
255 /**
256  * Deletes a (key, rid) pair from a B+ tree.
257  * <p>
258  * BPlusTree tree = new BPlusTree("t.txt", Type.intType(), 4);
259  * DataBox key = new IntDataBox(42);
260  * RecordId rid = new RecordId(42, (short) 42);
261  * <p>
262  * tree.put(key, rid);
263  * tree.get(key); // Optional.of(rid)
264  * tree.remove(key);
265  * tree.get(key); // Optional.empty()
266  */
267 public void remove(DataBox key) {
268     typecheck(key);
269     throw new UnsupportedOperationException("TODO: implement");
270 }
271
272 // Helpers //////////////////////////////////////
273
274 /**
275  * Returns a sexp representation of this tree. See BPlusNode.toSexp for
276  * more information.
277  */
278 public String toSexp() {
279     return root.toSexp();
280 }
281
282 /**
283  * Debugging large B+ trees is hard. To make it a bit easier, we can print
284  * out a B+ tree as a DOT file which we can then convert into a nice
285  * picture of the B+ tree. tree.toDot() returns the contents of DOT file
286  * which illustrates the B+ tree. The details of the file itself is not at
287  * all important, just know that if you call tree.toDot() and save the
288  * output to a file called tree.dot, then you can run this command
289  * <p>
290  * dot -T pdf tree.dot -o tree.pdf
291  * <p>
292  * to create a PDF of the tree.
293  */
294 public String toDot() {
295     List<String> strings = new ArrayList<>();
296     strings.add("digraph g {");
297     strings.add("    node [shape=record, height=0.1];");
298     strings.add(root.toDot());
299     strings.add("}");
300     return String.join("\n", strings);
301 }
302
303 /**
304  * Returns the largest number d such that the serialization of a LeafNode
305  * with 2d entries and an InnerNode with 2d keys will fit on a single page
306  * of size 'pageSizeInBytes'.
307  */
308 public static int maxOrder(int pageSizeInBytes, Type keySchema) {
309     int leafOrder = LeafNode.maxOrder(pageSizeInBytes, keySchema);
310     int innerOrder = InnerNode.maxOrder(pageSizeInBytes, keySchema);
311     return Math.min(leafOrder, innerOrder);
312 }
313
314 /**
315  * Returns the number of pages used to serialize the tree.
316  */
317 public int getNumPages() {
318     return metadata.getAllocator().getNumPages();
319 }
320

```

```

121  /**
122   * Serializes the header page to buf.
123   */
124  private void writeHeader(ByteBuffer buf) {
125      buf.put(metadata.getKeySchema().toBytes());
126      buf.putInt(metadata.getOrder());
127      buf.putInt(root.getPage().getPageNum());
128  }
129
130  private void typecheck(DataBox key) {
131      Type t = metadata.getKeySchema();
132      if (!key.type().equals(t)) {
133          String msg = String.format("DataBox %s is not of type %s", key, t);
134          throw new IllegalArgumentException(msg);
135      }
136  }
137
138  // Iterator //////////////////////////////////////
139  private class BPlusTreeIterator implements Iterator<RecordId> {
140      // A BPlusTreeIterator iterates over the entries of a B+ tree leaf by
141      // leaf. We maintain the following invariants:
142      //
143      // - leaf is null if and only if iter is null
144      // - iter is not null if and only if iter.hasNext()
145      private LeafNode leaf;
146      private Iterator<RecordId> iter;
147
148      public BPlusTreeIterator(LeafNode leaf, Iterator<RecordId> iter) {
149          assert (leaf != null);
150          assert (iter != null);
151          this.leaf = leaf;
152          this.iter = iter;
153
154          if (!this.iter.hasNext()) {
155              advance();
156          }
157      }
158
159      private void advance() {
160          Optional<LeafNode> sibling = leaf.getRightSibling();
161          if (sibling.isPresent()) {
162              this.leaf = sibling.get();
163              this.iter = this.leaf.scanAll();
164              if (!this.iter.hasNext()) {
165                  advance();
166              }
167          } else {
168              this.leaf = null;
169              this.iter = null;
170          }
171      }
172
173      @Override
174      public boolean hasNext() {
175          return iter != null;
176      }
177
178      @Override
179      public RecordId next() {
180          if (!hasNext()) {
181              throw new NoSuchElementException();
182          }
183          assert (leaf != null);
184          assert (iter != null);
185          assert (iter.hasNext());
186
187          RecordId rid = iter.next();
188          if (!iter.hasNext()) {
189              advance();
190          }
191      }
192  }

```

```

    }
    return rid;
}
}
}

```