


Chapter 11 :: Functional Languages

Programming Language Pragmatics

Michael L. Scott



Historical Origins

- The imperative and functional models grew out of work undertaken Alan Turing, Alonzo Church, Stephen Kleene, Emil Post, etc. ~1930s
 - different formalizations of the notion of an algorithm, or *effective procedure*, based on automata, symbolic manipulation, recursive function definitions, and combinatorics
- These results led Church to conjecture that *any* intuitively appealing model of computing would be equally powerful as well
 - this conjecture is known as *Church's thesis*

Historical Origins

- Turing's model of computing was the *Turing machine* a sort of pushdown automaton using an unbounded storage “tape”
 - the Turing machine computes in an imperative way, by changing the values in cells of its tape – like variables just as a high level imperative program computes by changing the values of variables

Historical Origins

- Church's model of computing is called the *lambda calculus*
 - based on the notion of parameterized expressions (with each parameter introduced by an occurrence of the letter λ —hence the notation's name.
 - Lambda calculus was the inspiration for functional programming
 - one uses it to compute by substituting parameters into expressions, just as one computes in a high level functional program by passing arguments to functions

Historical Origins

- Mathematicians established a distinction between
 - *constructive* proof (one that shows how to obtain a mathematical object with some desired property)
 - *nonconstructive* proof (one that merely shows that such an object must exist, e.g., by contradiction)
- Logic programming is tied to the notion of constructive proofs, but at a more abstract level:
 - the logic programmer writes a set of *axioms* that allow the *computer* to discover a constructive proof for each particular set of inputs

Functional Programming Concepts

- Functional languages such as Lisp, Scheme, FP, ML, Miranda, and Haskell are an attempt to realize Church's lambda calculus in practical form as a programming language
- The key idea: do everything by composing functions
 - no mutable state
 - no side effects

Functional Programming Concepts

- Necessary features, many of which are missing in some imperative languages
 - 1st class and high-order functions
 - serious polymorphism
 - powerful list facilities
 - structured function returns
 - fully general aggregates
 - garbage collection

Functional Programming Concepts

- So how do you get anything done in a functional language?
 - Recursion (especially tail recursion) takes the place of iteration
 - In general, you can get the effect of a series of assignments

```
x := 0      ...  
x := expr1  ...  
x := expr2  ...
```

from $f_3(f_2(f_1(0)))$, where each f expects the value of x as an argument, f_1 returns expr_1 , and f_2 returns expr_2

Functional Programming Concepts

- Recursion even does a nifty job of replacing looping

```
x := 0; i := 1; j := 100;
while i < j do
    x := x + i*j; i := i + 1;
    j := j - 1
end while
return x
```

becomes $f(0,1,100)$, where
 $f(x,i,j) ==$ if $i < j$ then
 $f(x+i*j, i+1, j-1)$ else x

Functional Programming Concepts

- Thinking about recursion as a direct, mechanical replacement for iteration, however, is the wrong way to look at things
 - One has to get used to thinking in a recursive style
- Even more important than recursion is the notion of *higher-order functions*
 - Take a function as argument, or return a function as a result
 - Great for building things

Functional Programming Concepts

- Lisp also has (these are not necessary present in other functional languages)
 - homo-iconography
 - self-definition
 - read-evaluate-print
- Variants of LISP
 - Pure (original) Lisp
 - Interlisp, MacLisp, Emacs Lisp
 - Common Lisp
 - Scheme

Functional Programming Concepts

- Pure Lisp is purely functional; all other Lisps have imperative features
- All early Lisps dynamically scoped
 - Not clear whether this was deliberate or if it happened by accident
- Scheme and Common Lisp statically scoped
 - Common Lisp provides dynamic scope as an option for explicitly-declared *special* functions
 - Common Lisp now THE standard Lisp
 - Very big; complicated (The Ada of functional programming)

Functional Programming Concepts

- Scheme is a particularly elegant Lisp
- Other functional languages
 - ML
 - Miranda
 - Haskell
 - FP
- Haskell is the leading language for research in functional programming

A Bit of Scheme

- As mentioned earlier, Scheme is a particularly elegant Lisp
 - Interpreter runs a read-eval-print loop
 - Things typed into the interpreter are evaluated (recursively) once
 - Anything in parentheses is a function call (unless quoted)
 - Parentheses are NOT just grouping, as they are in Algol-family languages
 - Adding a level of parentheses changes meaning
$$(+ \ 3 \ 4) \Rightarrow 7$$
$$((+ \ 3 \ 4)) \Rightarrow \text{error}$$
(the ' \Rightarrow ' arrow means 'evaluates to')

A Bit of Scheme

- Scheme:
 - Boolean values #t and #f
 - Numbers
 - Lambda expressions
 - Quoting
 - $(+ \ 3 \ 4) \Rightarrow 7$
 - $(\text{quote } (+ \ 3 \ 4)) \Rightarrow (+ \ 3 \ 4)$
 - $'(+ \ 3 \ 4) \Rightarrow (+ \ 3 \ 4)$
 - Mechanisms for creating new scopes
 - $(\text{let } ((\text{square } (\text{lambda } (x) (* \ x \ x))) (\text{plus } +))$
 $(\text{sqrt } (\text{plus } (\text{square } a) (\text{square } b))))$
 - let*
 - letrec

A Bit of Scheme

- Scheme:
 - Conditional expressions
`(if (< 2 3) 4 5) ⇒ 4`
`(cond`
 `((< 3 2) 1)`
 `((< 4 3) 2)`
 `(else 3)) ⇒ 3`
 - Imperative stuff
 - assignments
 - sequencing (`begin`)
 - iteration
 - I/O (`read`, `display`)

A Bit of Scheme

- Scheme standard functions (this is not a complete list):
 - arithmetic
 - boolean operators
 - equivalence
 - list operators
 - symbol?
 - number?
 - complex?
 - real?
 - rational?
 - integer?

A Bit of Scheme

Example program - Simulation of DFA

- We'll invoke the program by calling a function called 'simulate', passing it a DFA description and an input string
 - The automaton description is a list of three items:
 - start state
 - the transition function
 - the set of final states
 - The transition function is a list of pairs
 - the first element of each pair is a pair, whose first element is a state and whose second element is an input symbol
 - if the current state and next input symbol match the first element of a pair, then the finite automaton enters the state given by the second element of the pair

A Bit of Scheme

Example program - Simulation of DFA

```
(define simulate
  (lambda (dfa input)
    (letrec ((helper ; note that helper is tail recursive,
                  ; but builds the list of moves in reverse order
              (lambda (moves d2 i)
                (let ((c (current-state d2)))
                  (if (null? i) (cons c moves)
                      (helper (cons c moves) (move d2 (car i)) (cdr i))))))
      (let ((moves (helper '() dfa input)))
        (reverse (cons (if (is-final? (car moves) dfa)
                           'accept 'reject) moves)))))

;; access functions for machine description:
(define current-state car)
(define transition-function cadr)
(define final-states caddr)
(define is-final? (lambda (s dfa) (memq s (final-states dfa))))

(define move
  (lambda (dfa symbol)
    (let ((cs (current-state dfa)) (trans (transition-function dfa)))
      (list
        (if (eq? cs 'error)
            'error
            (let ((pair (assoc (list cs symbol) trans)))
              (if pair (cadr pair) 'error))) ; new start state
        trans ; same transition function
        (final-states dfa)))) ; same final states
```


Figure 11.1 Scheme program to simulate the actions of a DFA. Given a machine description and an input symbol i , function `move` searches for a transition labeled i from the start state to some new state s . It then returns a new machine with the same transition function and final states, but with s as its "start" state. The main function, `simulate`, encapsulates a tail-recursive helper function that accumulates an inverted list of moves, returning when it has consumed all input symbols. The wrapper then checks to see if the helper ended in a final state; it returns the (properly ordered) series of moves, with `accept` or `reject` at the end. The functions `cadr` and `caddr` are defined as `(lambda (x) (car (cdr x)))` and `(lambda (x) (car (cdr (cdr x))))`, respectively. Scheme provides a large collection of such abbreviations.



A Bit of Scheme

Example program - Simulation of DFA

Start



```
(define zero-one-even-dfa
```

```
  '(q0                                     ; start state
    (((q0 0) q2) ((q0 1) q1) ((q1 0) q3) ((q1 1) q0) ; transition fn
     ((q2 0) q0) ((q2 1) q3) ((q3 0) q1) ((q3 1) q2))
    (q0)))                                     ; final states
```

Figure 11.2 DFA to accept all strings of zeros and ones containing an even number of each. At the bottom of the figure is a representation of the machine as a Scheme data structure, using the conventions of Figure 11.1.

A Bit of OCaml

- OCaml is a descendent of ML, and cousin to Haskell, F#
 - “O” stands for objective, referencing the object orientation introduced in the 1990s
 - Interpreter runs a read-eval-print loop like in Scheme
 - Things typed into the interpreter are evaluated (recursively) once
 - Parentheses are NOT function calls, but indicate tuples

A Bit of OCaml

- Ocaml:
 - Boolean values
 - Numbers
 - Chars
 - Strings
 - More complex types created by lists, arrays, records, objects, etc.
 - $(+ \ - \ * \ /)$ for ints, $(+.\ -. \ *.\ ./.)$ for floats
 - `let` keyword for creating new names
 - `let average = fun x y -> (x +. y) /. 2.;;`

A Bit of OCaml

- Ocaml:

- Variant Types

```
type 'a tree = Empty | Node of 'a * 'a tree * 'a tree;;
```

- Pattern matching

```
let atomic_number (s, n, w) = n;;  
let mercury = ("Hg", 80, 200.592);;  
atomic_number mercury;;    ⇒ 80
```

A Bit of OCaml

- OCaml:
 - Different assignments for references ``:='` and array elements ``<-'`

```
let insertion_sort a =  
  for i = 1 to Array.length a - 1 do  
    let t = a.(i) in  
    let j = ref i in  
    while !j > 0 && t < a.(!j - 1) do  
      a.(!j) <- a.(!j - 1);  
      j := !j - 1  
    done;  
    a.(!j) <- t  
  done;;
```


A Bit of OCaml

Example program - Simulation of DFA

- We'll invoke the program by calling a function called 'simulate', passing it a DFA description and an input string
 - The automaton description is a record with three fields:
 - start state
 - the transition function
 - the list of final states
 - The transition function is a list of triples
 - the first two elements are a state and an input symbol
 - if these match the current state and next input, then the automaton enters a state given by the third element

A Bit of OCaml

Example program - Simulation of DFA

```
open List;;      (* includes rev, find, and mem functions *)

type state = int;;
type 'a dfa = {
  current_state : state;
  transition_function : (state * 'a * state) list;
  final_states : state list;
};;
type decision = Accept | Reject;;

let move (d:'a dfa) (x:'a) : 'a dfa =
  { current_state = (
    let (_, _, q) =
      find (fun (s, c, _) -> s = d.current_state && c = x)
          d.transition_function in
    q);
    transition_function = d.transition_function;
    final_states = d.final_states;
  };;

let simulate (d:'a dfa) (input:'a list) : (state list * decision) =
  let rec helper moves d2 remaining_input : (state option * state list) =
    match remaining_input with
    | [] -> (Some d2.current_state, moves)
    | hd :: tl ->
      let new_moves = d2.current_state :: moves in
      try helper new_moves (move d2 hd) tl
      with Not_found -> (None, new_moves) in
  match helper [] d input with
  | (None, moves) -> (rev moves, Reject)
  | (Some last_state, moves) ->
    ( rev (last_state :: moves),
      if mem last_state d.final_states then Accept else Reject);;
```


Figure 11.3 OCaml program to simulate the actions of a DFA. Given a machine description and an input symbol i , function `move` searches for a transition labeled i from the start state to some new state s . If the search fails, `find` raises exception `Not_found`, which propagates out of `move`; otherwise `move` returns a new machine with the same transition function and final states, but with s as its “start” state. Note that the code is polymorphic in the type of the input symbols. The main function, `simulate`, encapsulates a tail-recursive helper function that accumulates an inverted list of moves, returning when it has consumed all input symbols. The encapsulating function then checks to see if the helper ended in a final state; it returns the (properly ordered) series of moves, together with an `Accept` or `Reject` indication. The built-in option constructor (Example 7.6) is used to distinguish between a real state (`Some s`) and an error state (`None`).



A Bit of OCaml

Example program - Simulation of DFA

Start



```
let a_b_even_dfa : char dfa =
  { current_state = 0;
    transition_function =
      [ (0, 'a', 2); (0, 'b', 1); (1, 'a', 3); (1, 'b', 0);
        (2, 'a', 0); (2, 'b', 3); (3, 'a', 1); (3, 'b', 2) ];
    final_states = [0];
  };
```

Figure 11.4 DFA to accept all strings of as and bs containing an even number of each. At the bottom of the figure is a representation of the machine as an OCaml data structure, using the conventions of Figure 11.3.

Evaluation Order Revisited

- Applicative order
 - what you're used to in imperative languages
 - usually faster
- Normal order
 - like call-by-name: don't evaluate arg until you need it
 - sometimes faster
 - terminates if anything will (Church-Rosser theorem)

Evaluation Order Revisited

- In Scheme
 - functions use applicative order defined with lambda
 - special forms (aka macros) use normal order defined with syntax-rules
- A *strict* language requires all arguments to be well-defined, so applicative order can be used
- A *non-strict* language does not require all arguments to be well-defined; it requires normal-order evaluation

Evaluation Order Revisited

- Lazy evaluation gives the best of both worlds
- But not good in the presence of side effects.
 - delay and force in Scheme
 - delay creates a "promise"

High-Order Functions

- Higher-order functions
 - Take a function as argument, or return a function as a result
 - Great for building things
 - Currying (after Haskell Curry, the same guy Haskell is named after)
 - For details see Lambda calculus on CD
 - ML, Miranda, OCaml, and Haskell have especially nice syntax for curried functions

Functional Programming in Perspective

- Advantages of functional languages
 - lack of side effects makes programs easier to understand
 - lack of explicit evaluation order (in some languages) offers possibility of parallel evaluation (e.g. MultiLisp)
 - lack of side effects and explicit evaluation order simplifies some things for a compiler (provided you don't blow it in other ways)
 - programs are often surprisingly short
 - language can be extremely small and yet powerful



Functional Programming in Perspective

- Problems

- difficult (but not impossible!) to implement efficiently on von Neumann machines

- lots of copying of data through parameters
 - (apparent) need to create a whole new array in order to change one element
 - heavy use of pointers (space/time and locality problem)
 - frequent procedure calls
 - heavy space use for recursion
 - requires garbage collection
 - requires a different mode of thinking by the programmer
 - difficult to integrate I/O into purely functional model



ELSEVIER