

Programming Project 1

Instructions for this assignment:

You may work on this assignment either solo, or in a group of up to four students (including yourself). Submit the assignment by following the steps listed under “Submission Procedure” before the due date/time; late submissions will receive a grade of zero, but on-time submissions will receive partial credit even if incomplete. If working in a group, each member of the group should submit the assignment, and enter the names of all group members in the Comments field in Blackboard.

This project is intended to familiarize you with the Linux operating system, and introduce you to the elements of concurrent programming using the POSIX standard API. Most Unix and Linux systems adhere closely to this API, and these or similar systems are heavily used in backend applications and cloud computing.

Rather than using a Linux system on a real computer, you will be using an Ubuntu Linux Virtual Machine (provided by the authors of our required textbook) running on the VirtualBox application. This will allow you to run the Ubuntu operating system on your own personal computer without any system changes. Using the Ubuntu virtual machine allows you administrator access without the risk of modifying your host operating system, permitting access to operating system functionality.

This project as a whole (Programming Project 1 & 2) should be viewed as a virtual lab. It is intended to allow you to learn how to use the Linux operating system and program on the system. As such, it consists of a number of smaller steps, leading to the development of some simple concurrent programs and installation of simple kernel modules.

Step 1

The first step of the project is simply to create and begin using your Ubuntu Linux virtual machine. You should follow the instructions at the website provided by textbook authors Silberschatz, Galvin, and Gagne [1] for downloading and installing both VirtualBox and the Ubuntu Virtual Disk Image on your own computer, and to create the virtual machine. You can even create and run the Linux virtual machine by copying your VDI onto a thumb drive, and carry it with you.

[1] <http://people.westminstercollege.edu/faculty/ggagne/osc10e/vm/index.html>

Once you have created the virtual machine, start it and open a terminal window. In your terminal window, do the following:

1. Type

```
touch <your first name>
```

and press the enter key, where **<your first name>** is replaced with your actual first name.

(Note: You cannot have spaces in the name that you give the `touch` command.) This will create a Linux file with your name as the filename.

2. Type `ls` and press the enter key to display the files in your current directory.
3. Type `date` and press the enter key to display the current date and time.
4. **Take a screenshot of your virtual machine showing the terminal window after running the `date` command. Save your screenshot as a GIF image with the filename `ScreenShot-1.gif`. (You will submit this screenshot as your deliverable for Step 1 when done with all steps.)**

Step 2

In this step, you will modify one of the sample programs using the `fork` and `exec` system calls to display your name before displaying the files in the current folder.

To begin, open a terminal window and type:

```
mkdir Project1Step2
```

and press the enter key to create a new folder called `Project1Step2`. Now type:

```
cd Project1Step2
```

and press the enter key to make `Project1Step2` the new current working directory.

Type:

```
pwd
```

and press the enter key to display the pathname of the new working directory.

Now type:

```
cp ~/COMP3080/final-src-osc9e/ch3/newproc-posix.c .
```

and press the enter key, to copy the file `new-proc.c` from the OSC textbook source code directory to your new project folder.

Next, type:

```
gedit newproc-posix.c &
```

and press the enter key to open the program file `newproc-posix.c` in the `gedit` editor program (or use your favorite editor). Modify line 27 of this file, which reads:

```
printf("I am the child %d\n",pid);
```

to read as follows:

```
printf("I am <your-name> in process %d\n", pid);
```

where again you replace `<your-name>` with your actual name.

In the editor, save the file (if using `gedit`, go to the menu and click `File->Save` or simply press `Ctl-s`).

In the terminal, type:

```
gcc -o newproc-posix newproc-posix.c
```

and press enter to compile your modified program.

Type:

```
./newproc-posix
```

and press enter to run your modified program.

Take a screenshot of your virtual machine's terminal window, and name it `ScreenShot-2.gif`.

Step 3

For this part of the project, we'll explore an interesting part of the Linux file system: The `/proc` directory (folder). The `/proc` directory holds a set of files and directories which represent operating system information. For example, through the `/proc` directory, you can find information about your system's CPU, devices, memory, etc. Additionally, every process running on a Linux system is represented in the `/proc` directory by a directory with the process PID.

Open a terminal window and type:

```
cd /proc
```

and press the enter key (this will be the last time I tell you to press the enter key after typing a command – yay!) to use `/proc` as the current working directory.

Open a second terminal window, and type:

```
cd /proc
```

In one of the windows, type:

```
ps
```

You will see a display that looks something like the following:

PID	TTY	TIME	CMD
17851	pts/1	00:00:00	bash
17875	pts/1	00:00:00	ps

The number underneath the PID column on the line ending with the name `bash` is the process ID for the terminal window in which you typed the command "`ps`".

In the other terminal window, type:

```
ls
```

and look for a directory with the process ID you just found as its name. Type:

```
cd <processID>
```

where `<processID>` is the process ID you found, to enter that directory.

Then type:

```
cd fd
```

to enter the `fd` directory. Type:

```
ls
```

You will see four directories with the names: `0`, `1`, `2`, `255`. These directories represent file descriptors for files opened by the other terminal window. `0` is the standard input (`stdin`), by default tied to the keyboard; `1` is the standard output (`stdout`), by default tied to the monitor; and `2` is the standard error output (`stderr`), by default tied to the monitor. `255` is a special file used by the terminal window and its shell internally.

The Linux `cat` command echoes whatever is typed on the keyboard to the default display. However, in Linux, you can redirect the input and output of a command like `cat` so that the command takes its input from and sends its output to any other source you like. The character `<` is used to redirect another source (usually a file) to the default input of `cat`, and `>` is used to redirect the default output of `cat` to another source (again, usually a file).

In the terminal window in the `fd` directory, type:

```
cat >2
```

and watch the other terminal window as you type:

```
Hello from <your Name>
```

replacing `<your Name>` with your own name. Type `Ctrl-d` (by pressing the Control or `Ctrl` key at the same time as the `d` key) to close the `cat` command. (Note: The `Ctrl-d` character is the Linux end-of-file character.)

Based on what you observed, and the redirection used in the `cat >2` command, what is your interpretation of what happened? Write up your answer in a file called `Step3-README.pdf` for later submission (see below).

Also, take a screenshot showing both of your terminal windows after the `cat >2` command was executed, and name it `ScreenShot-3.gif`.

Step 4

In this part of the project, you are going to redirect Linux input within a C program after forking a child process and using the `execlp` system call. I'll first provide some background.

The Linux/UNIX shell command line environment provides a large number of simple commands that process text files. Suppose that you have a text file, `Names.txt`, that contains a series of names (one per line of the file), and that you want to know how many names are in the file. You can type:

```
wc -l < Names.txt
```

which will display the number of lines in `Names.txt`, which is the number of names.

However, if you examine `Names.txt`, you'll see that it contains a lot of duplicate names. If you want to know the number of unique names in `Names.txt`, you can use the Linux shell's pipe redirector to use several commands in a pipeline, as below:

```
sort < Names.txt | uniq | wc -l
```

The `sort` command is given the list of names and sorts them. The `uniq` command removes duplicate lines in a file, but only if the duplicate lines are adjacent. By sorting `Names.txt` and piping the sorted list to `uniq` (using the `"|"` symbol), the duplicate lines are adjacent and `uniq` can remove them. Finally, the sorted list with duplicates removed is piped to the `wc` command, which counts the number of lines of the sorted list with the duplicates removed to display the number of unique names.

You will write a small C program which does what the Linux shell does, without doing the command line processing. That is, you'll write a program that creates three child processes, one for each command. Each child process will tie an end of a pipe to the standard input, standard output, or both, and then call the `execlp` system call to run the command. The three child processes will be connected by two Linux pipes, one pipe between the first two processes and the other pipe between the second two processes. The overall structure of the program is shown in Figure 1. Figures 2 and 3 show examples used in class which show the use of `execlp` to run a Linux command and the creation of pipes. Your program will use the techniques from both of these programs.

However, the code in Figure 3, which demonstrates pipes, doesn't handle the redirection needed. The problem is that once the `execlp` finishes, each child process is running a brand-new program image which knows nothing of the pipes that were created by the parent process. So, before calling `execlp` to start the new program, each child must first tie the appropriate ends of the pipes to its standard input and/or output using the `dup2` system call.

Your program will need to use the following include statements:

```
#include <stdio.h>
#include <unistd.h>
```

You will also need the following Linux system calls:

- `printf(<format-string>, arg0, arg1, ..., argn);`
`fprintf(<file-pointer>, <format-string>, arg0, arg1, ..., argn);`

`printf` and `fprintf` allow you to display formatted output. The only difference between them is that `printf` always displays its output on the standard output (`stdout`), while `fprintf` allows you to choose the output file (for this assignment, the only need for `fprintf` is to send output to the standard error output, `stderr`). Replace `<file-pointer>` with the file pointer to which you want to send the output. The two standard output file pointers are `stdout` and `stderr`. The `<format-string>` is a C-language literal string that contains the text to send to output with embedded placeholders (`%d`, integer value, `%s`, string value, `%f`, floating point value, etc.). Each placeholder in the format string is replaced with one of the remaining arguments (`arg0`, `arg1`, etc.) in the call. There should be as many arguments as there are placeholders in the format string. The call:

```
printf("This process id is %d\n", getpid());
```

displays the string:

```
This process id is 19334
```

on a separate line, if the function `getpid()` returns 19334 as the process ID for the current process.

- `getpid()`
This system call returns the process ID (pid) of the current process.
- `fork()`
This system call creates a new process that is an exact duplicate of the current process. The only difference between the two processes are the return values from `fork()`: The child process sees a return value of zero, while the parent process sees the PID of the child process.
- `close(<fd>)`
This system call closes the file descriptor `<fd>` for output. You do not need to check for errors.
- `pipe(<fd-array>)`
This system call places two file descriptors into the two-element integer array `<fd-array>`

corresponding to the ends of an open pipe. Element zero of `<fd-array>` is the read end of the pipe, and element one is the write end of the pipe.

- `dup2(<oldfd>, <newfd>)`

This system call ties the existing file descriptor `<oldfd>` to the file descriptor `<newfd>`. After the call, `<oldfd>` and `<newfd>` refer to the same file descriptor. If `<newfd>` is already open, the current output is closed and then `<newfd>` is tied to `<oldfd>`. For this part of the project, `<oldfd>` will be either the read or write end of one of the pipes (`fd1` or `fd2`), and `<newfd>` will be either 0 (standard input) or 1 (standard output). You do not need to check for errors.

- `execvp(<command-path-name>, <command-name>, <arg-string-1>, <arg-string-2>, ..., NULL);`

This system call replaces the current program image (code, data, and stack) with the image found at the path name `<command-path-name>`. After this call is complete, the process is running a completely different program with new code, data, and stack. `<command-name>` is the simple name for the command, which is placed in the command line argument zero. `<arg-string-1>`, `<arg-string-2>`, etc. are command line arguments 1, 2, etc. The command line argument strings are C-language pointers to character arrays (as are C-language string literals). The set of command line arguments must end with a zero or `NULL` pointer. If there are no arguments for the command, simply use `NULL`. The parameters to `execvp` for each of the commands are as follows:

- `sort`

- `<command-path-name>` is `"/usr/bin/sort"`
- `<command-name>` is `"sort"`
- no `<arg-string-x>`, only `NULL`

- `uniq`

- `<command-path-name>` is `"/usr/bin/uniq"`
- `<command-name>` is `"uniq"`
- no `<arg-string-x>`, only `NULL`

- `wc`

- `<command-path-name>` is `"/usr/bin/wc"`
- `<command-name>` is `"wc"`
- `<arg-string-1>` is `"-l"`
- `NULL`

- `wait(<pid>, NULL)`

This system call waits for the child process with the process id `<pid>`. The `NULL` second argument indicates that no status information is to be returned.

We're now ready to begin work on Step 4.

1. To start this step of the project, open a terminal window and type:

```
mkdir SortUniqWc
```

and then:

```
cd SortUniqWc
```

2. You will now be in the directory `SortUniqWc`. Type:

```
gedit sortuniqwc.c &
```

to open the `gedit` editor (or open your favorite editor). Use this editor to write your program. Use the output provided in Figure 1 and the examples shown in Figures 2 and 3 to guide you.

It's a good idea to add `printf` statements to each of the children to display their process IDs:

```
printf("The child process running <cmd> is %d\n",  
      getpid());
```

where `<cmd>` is the name of the command this child is running.

Also add a `printf` after each `execlp`:

```
printf("Should not be here after execlp to <cmd>\n");
```

If `execlp` fails, the `printf` will display a message; if the `execlp` succeeds, the `printf` will not execute. This is a simple way to verify that `execlp` is working. Typical reasons for `execlp` failing are a bad command path name, or a bad argument sequence (most commonly, a missing `NULL` at the end of the command line arguments).

Also, add `printf` to the main process code that informs you when the last child is finished (just after the call to `wait`).

3. From the Blackboard page for this assignment, download the file, `Names.txt`, that contains the list of names. Import this file into your Ubuntu VM either by using the shared folder facility, or by setting the virtual machine to accept copy and paste from the host OS. For the latter, click on VirtualBox's `Devices` menu and select `Drag and Drop->Bidirectional`.

4. Type the command:

```
sort < Names.txt | uniq | wc -l
```

in your terminal. You should get a display of 23, indicating that there are 23 unique names in `Names.txt`.

5. To run your program, you must first compile it. To do so, type

```
gcc -o sortuniqwc sortuniqwc.c
```

This command will compile the program in `sortuniqwc.c` and put the binary executable in the file `sortuniqwc`.

6. To run your program, type:

```
./sortuniqwc < Names.txt
```

The `./` tells the shell to execute the program in the file `sortuniqwc` by looking in the current directory (`SortUniqWd`) rather than searching through the standard system paths. You should see the printed messages you added to the code. There is no certain order to these messages; running the program several times may result in slightly different orders. This is normal.

Submission Procedure

Gather together the following files into a single directory with the name “<YourName>-P1” (e.g., “TomWilkes-P1”):

ScreenShot-1.gif, ScreenShot-2.gif, ScreenShot-3.gif,
Step3-README.pdf, and sortuniqwc.c

Now, use the `tar` and `gzip` commands to compress your directory containing these files into a single file (known as a “tarball”) called “<YourName>-P1.tar.gz” (e.g., “TomWilkes-P1.tar.gz”). For example, assuming that the current working directory is the directory that contains TomWilkes-P1:

```
tar cvf TomWilkes-P1.tar TomWilkes-P1
gzip TomWilkes.tar
```

Finally, submit your `.tar.gz` file using the Blackboard page for this assignment.

Alternatively, if your Ubuntu system supports creating a Zipped folder directly, you may submit a `.zip` file of your project folder, using the same naming scheme (e.g., “TomWilkes-P1.zip”).

Remember to indicate in the Comments field for the submission whether you worked solo or in a group, and (if you worked in a group) list all of your group members including yourself.

Outline for your Step 4 Program

```
/*  include files here  */

int main(int argc, char *argv[]) {

    //create first pipe fd1
    // fork first child
    pid = fork(); // create first child for sort
    if (pid < 0) {
        // fork error
    }
    if (pid == 0) { // first child process, run sort
        // tie write end of pipe fd1 to standard output (file descriptor
        1)
        // close read end of pipe fd1
        // start the sort command using execlp
        // should not get here
    }
    //create second pipe fd2
    // fork second child
    pid = fork(); // create second child for uniq
    if (pid < 0) {
        // fork error
    }
    if (pid == 0) { // second child process, run uniq
        // tie read end of fd1 to standard input (file descriptor 0)
        // tie write end of fd2 to standard output (file descriptor 1)
        // close write end of pipe fd1
        // close read end of pipe fd2
        // start the uniq command using execlp
        // should not get here
    }
    // fork third child
    pid = fork() // create third child for wc -l
    if (pid < 0) {
        // fork error
    }
    if (pid == 0) { // third child process, run wc -l
        // tie read end of fd2 to standard input (file descriptor 0)
        // close write end of pipe fd2
        // close read end of pipe fd1
        // close write end of pipe fd1
        // start the wc -l command using execlp
        // should not get here
    }
    // parent process code
    // close both ends of pipes fd1 and fd2
    // wait for third process to end.
}
```

Figure 1: Outline of the *sortuniqwc.c* program

Example of using `exec1p`

```
/**
 * This program forks a separate process using the fork()/exec() system calls.
 *
 * Figure 3.09
 *
 * @author Silberschatz, Galvin, and Gagne
 * Operating System Concepts - Ninth Edition
 * Copyright John Wiley & Sons - 2013
 */

#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("I am the child %d\n", pid);
        exec1p("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        printf("I am the parent %d\n", pid);
        wait(NULL);

        printf("Child Complete\n");
    }

    return 0;
}
```

Figure 2: Figure 3.9 from textbook showing use of `exec1p` to run the Linux command `ls`

Example using the pipe system call

```
/**
 * Example program demonstrating UNIX pipes.
 *
 * Figures 3.25 & 3.26
 *
 * @author Silberschatz, Galvin, and Gagne
 * Operating System Concepts - Ninth Edition
 * Copyright John Wiley & Sons - 2013
 */

#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <string.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    pid_t pid;
    int fd[2];

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    /* now fork a child process */
    pid = fork();

    if (pid < 0) {
        fprintf(stderr, "Fork failed");
        return 1;
    }

    if (pid > 0) { /* parent process */
        /* close the unused end of the pipe */
        close(fd[READ_END]);

        /* write to the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

        /* close the write end of the pipe */
        close(fd[WRITE_END]);
    }
    else { /* child process */
        /* close the unused end of the pipe */
        close(fd[WRITE_END]);

        /* read from the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE);
        printf("child read %s\n", read_msg);

        /* close the write end of the pipe */
        close(fd[READ_END]);
    }

    return 0;
}
```

Figure 3: Figures 3.25 and 3.26 from textbook showing use of Linux pipes